



Node servers

ON THIS PAGE



To generate a standalone Node server, use [adapter-node](#) .

Usage

Install with `npm i -D @sveltejs/adapter-node` , then add the adapter to your `svelte.config.js` :

svelte.config.js

```
import adapter from '@sveltejs/adapter-node';

export default {
  kit: {
    adapter: adapter()
  }
};
```

Deploying

First, build your app with `npm run build` . This will create the production server in the output directory specified in the adapter options, defaulting to `build` .

You will need the output directory, the project's `package.json` , and the production dependencies in `node_modules` to run the application. Production dependencies can be generated by copying the `package.json` and `package-lock.json` and then running `npm ci`



Development dependencies will be bundled into your app using Rollup. To control whether a given package is bundled or externalised, place it in `devDependencies` or `dependencies` respectively in your `package.json`.

Compressing responses

You will typically want to compress responses coming from the server. If you are already deploying your server behind a reverse proxy for SSL or load balancing, it typically results in better performance to also handle compression at that layer since Node.js is single-threaded.

However, if you're building a custom server and do want to add a compression middleware there, note that we would recommend using @polka/compression since SvelteKit streams responses and the more popular `compression` package does not support streaming and may cause errors when used.

Environment variables

In `dev` and `preview`, SvelteKit will read environment variables from your `.env` file (or `.env.local`, or `.env.[mode]`, as determined by Vite.)

In production, `.env` files are *not* automatically loaded. To do so, install `dotenv` in your project...

```
npm install dotenv
```

...and invoke it before running the built app:

```
node -r dotenv/config build
```

If you use Node.js v20.6+, you can use the --env-file flag instead:

PORT, HOST and SOCKET_PATH

By default, the server will accept connections on `0.0.0.0` using port 3000. These can be customised with the `PORT` and `HOST` environment variables:

```
HOST=127.0.0.1 PORT=4000 node build
```

Alternatively, the server can be configured to accept connections on a specified socket path. When this is done using the `SOCKET_PATH` environment variable, the `HOST` and `PORT` environment variables will be disregarded.

```
SOCKET_PATH=/tmp/socket node build
```

ORIGIN, PROTOCOL_HEADER, HOST_HEADER, and PORT_HEADER

HTTP doesn't give SvelteKit a reliable way to know the URL that is currently being requested. The simplest way to tell SvelteKit where the app is being served is to set the `ORIGIN` environment variable:

```
ORIGIN=https://my.site node build

# or e.g. for local previewing and testing
ORIGIN=http://localhost:3000 node build
```

With this, a request for the `/stuff` pathname will correctly resolve to `https://my.site/stuff`. Alternatively, you can specify headers that tell SvelteKit about the request protocol and host, from which it can construct the origin URL:

```
PROTOCOL_HEADER=x-forwarded-proto HOST_HEADER=x-forwarded-host node build
```

original protocol and host if you're using a reverse proxy (think load balancers and CDNs). You should only set these variables if your server is behind a trusted reverse proxy; otherwise, it'd be possible for clients to spoof these headers.

If you're hosting your proxy on a non-standard port and your reverse proxy supports `x-forwarded-port`, you can also set `PORT_HEADER=x-forwarded-port`.

If `adapter-node` can't correctly determine the URL of your deployment, you may experience this error when using form actions:

Cross-site POST form submissions are forbidden

ADDRESS_HEADER and XFF_DEPTH

The `RequestEvent` object passed to hooks and endpoints includes an `event.getClientAddress()` function that returns the client's IP address. By default this is the connecting `remoteAddress`. If your server is behind one or more proxies (such as a load balancer), this value will contain the innermost proxy's IP address rather than the client's, so we need to specify an `ADDRESS_HEADER` to read the address from:

```
ADDRESS_HEADER=True-Client-IP node build
```

Headers can easily be spoofed. As with `PROTOCOL_HEADER` and `HOST_HEADER`, you should know what you're doing before setting these.

If the `ADDRESS_HEADER` is `X-Forwarded-For`, the header value will contain a comma-separated list of IP addresses. The `XFF_DEPTH` environment variable should specify how many trusted proxies sit in front of your server. E.g. if there are three trusted proxies, proxy 3 will forward the addresses of the original connection and the first two proxies:

```
<client address>, <proxy 1 address>, <proxy 2 address>
```

spoofing:

```
<spoofed address>, <client address>, <proxy 1 address>, <proxy 2 address>
```

We instead read from the *right*, accounting for the number of trusted proxies. In this case, we would use `XFF_DEPTH=3` .

If you need to read the left-most address instead (and don't care about spoofing) — for example, to offer a geolocation service, where it's more important for the IP address to be *real* than *trusted*, you can do so by inspecting the `x-forwarded-for` header within your app.

BODY_SIZE_LIMIT

The maximum request body size to accept in bytes including while streaming. The body size can also be specified with a unit suffix for kilobytes (`K`), megabytes (`M`), or gigabytes (`G`). For example, `512K` or `1M` . Defaults to `512kb`. You can disable this option with a value of `Infinity` (0 in older versions of the adapter) and implement a custom check in `handle` if you need something more advanced.

SHUTDOWN_TIMEOUT

The number of seconds to wait before forcefully closing any remaining connections after receiving a `SIGTERM` or `SIGINT` signal. Defaults to `30` . Internally the adapter calls `closeAllConnections` . See [Graceful shutdown](#) for more details.

IDLE_TIMEOUT

When using `systemd` socket activation, `IDLE_TIMEOUT` specifies the number of seconds after which the app is automatically put to sleep when receiving no requests. If not set, the app runs continuously. See [Socket activation](#) for more details.

The adapter can be configured with various options:

svelte.config.js

```
import adapter from '@sveltejs/adapter-node';

export default {
  kit: {
    adapter: adapter({
      // default options are shown
      out: 'build',
      precompress: true,
      envPrefix: ''
    })
  }
};
```

out

The directory to build the server to. It defaults to `build` — i.e. `node build` would start the server locally after it has been created.

precompress

Enables precompressing using gzip and brotli for assets and prerendered pages. It defaults to `true`.

envPrefix

If you need to change the name of the environment variables used to configure the deployment (for example, to deconflict with environment variables you don't control), you can specify a prefix:

```
envPrefix: 'MY_CUSTOM_';
```

```
MY_CUSTOM_PORT=4000 \
MY_CUSTOM_ORIGIN=https://my.site \
node build
```

Graceful shutdown

By default `adapter-node` gracefully shuts down the HTTP server when a `SIGTERM` or `SIGINT` signal is received. It will:

1. reject new requests (`server.close`)
2. wait for requests that have already been made but not received a response yet to finish and close connections once they become idle (`server.closeIdleConnections`)
3. and finally, close any remaining connections that are still active after `SHUTDOWN_TIMEOUT` seconds. (`server.closeAllConnections`)

If you want to customize this behaviour you can use a [custom server](#).

You can listen to the `sveltekit:shutdown` event which is emitted after the HTTP server has closed all connections. Unlike Node's `exit` event, the `sveltekit:shutdown` event supports asynchronous operations and is always emitted when all connections are closed even if the server has dangling work such as open database connections.

```
process.on('sveltekit:shutdown', async (reason) => {
  await jobs.stop();
  await db.close();
});
```

The parameter `reason` has one of the following values:

`SIGINT` - shutdown was triggered by a `SIGINT` signal

`SIGTERM` - shutdown was triggered by a `SIGTERM` signal

Most Linux operating systems today use a modern process manager called systemd to start the server and run and manage services. You can configure your server to allocate a socket and start and scale your app on demand. This is called socket activation. In this case, the OS will pass two environment variables to your app — `LISTEN_PID` and `LISTEN_FDS`. The adapter will then listen on file descriptor 3 which refers to a systemd socket unit that you will have to create.

You can still use `envPrefix` with systemd socket activation. `LISTEN_PID` and `LISTEN_FDS` are always read without a prefix.

To take advantage of socket activation follow these steps.

1. Run your app as a systemd service. It can either run directly on the host system or inside a container (using Docker or a systemd portable service for example). If you additionally pass an `IDLE_TIMEOUT` environment variable to your app it will gracefully shutdown if there are no requests for `IDLE_TIMEOUT` seconds. systemd will automatically start your app again when new requests are coming in.

```
/etc/systemd/system/myapp.service
```

```
[Service]
```

```
Environment=NODE_ENV=production IDLE_TIMEOUT=60
```

```
ExecStart=/usr/bin/node /usr/bin/myapp/build
```

2. Create an accompanying socket unit. The adapter only accepts a single socket.

```
/etc/systemd/system/myapp.socket
```

```
[Socket]
```

```
ListenStream=3000
```

```
[Install]
```

```
WantedBy=sockets.target
```


is made to `localhost:3000` .

Custom server

The adapter creates two files in your build directory — `index.js` and `handler.js` .

Running `index.js` — e.g. `node build` , if you use the default build directory — will start a server on the configured port.

Alternatively, you can import the `handler.js` file, which exports a handler suitable for use with Express, Connect or Polka (or even just the built-in `http.createServer`) and set up your own server:

my-server.js

```
import { handler } from './build/handler.js';
import express from 'express';

const app = express();

// add a route that lives separately from the SvelteKit app
app.get('/healthcheck', (req, res) => {

  res.end('ok');
});

// let SvelteKit handle everything else, including serving prerendered pages and static assets
app.use(handler);

app.listen(3000, () => {
  console.log('listening on port 3000');
});
```

[Edit this page on GitHub](#)

PREVIOUS

[Zero-config deployments](#)

NEXT

[Static site generation](#)

